

**Advanced DevOps Lab  
Experiment:3**

**Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.**

**Theory:**

To understand Kubernetes Cluster Architecture and how to install and spin up a Kubernetes cluster on Linux machines or cloud platforms, it's essential to grasp the fundamental components and design principles of Kubernetes.

**Overview of Kubernetes**

Kubernetes is an open-source container orchestration platform developed by Google, designed to automate the deployment, scaling, and management of containerized applications. It provides a robust infrastructure that supports microservices architecture, offering features such as self-healing, scaling, and zero-downtime deployments. Kubernetes can run on various environments, including public clouds (like AWS and Azure), private clouds, and bare metal servers.

**Kubernetes Architecture**

Kubernetes architecture is primarily composed of two main components: the Control Plane and the Data Plane.

**Control Plane**

The Control Plane manages the overall state of the Kubernetes cluster and includes several key components:

- kube-apiserver: The API server acts as the gateway for all interactions with the cluster, processing REST requests and managing the state of the cluster.
- etcd: A distributed key-value store that holds the configuration data and state of the cluster, ensuring consistency and availability.
- kube-scheduler: Responsible for assigning Pods to worker nodes based on resource availability and other constraints.
- kube-controller-manager: Manages controllers that regulate the state of the cluster, ensuring that the desired state matches the actual state.

- cloud-controller-manager (optional): Integrates with cloud provider APIs to manage resources specific to the cloud environment.

## Data Plane

The Data Plane consists of the worker nodes that run the containerized applications. Each worker node includes:

- kubelet: An agent that ensures containers are running in Pods. It communicates with the Control Plane to receive instructions.
- kube-proxy: Maintains network rules and facilitates communication between Pods and services.
- Container Runtime: Software responsible for running containers, such as Docker or containerd.

## Core Concepts

Key concepts in Kubernetes include:

- Pods: The smallest deployable units in Kubernetes, which can contain one or more containers.
- Services: Abstracts a set of Pods, providing a stable network endpoint for accessing them.
- Deployments: Define the desired state for Pods and manage their lifecycle, including scaling and updates.

## Installing and Spinning Up a Kubernetes Cluster

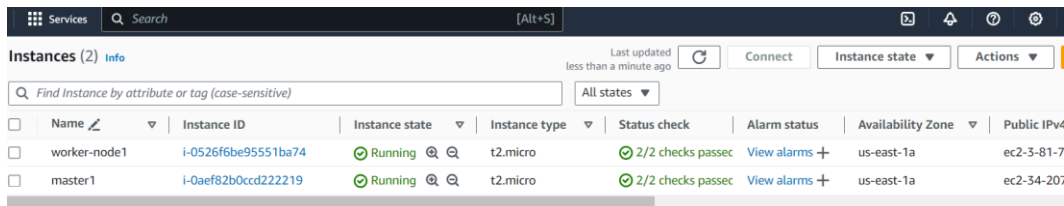
To install and set up a Kubernetes cluster, follow these general steps:

1. Choose an Environment: Decide whether to deploy on local machines or a cloud platform. For cloud platforms, services like Google Kubernetes Engine (GKE), Amazon EKS, or Azure AKS can simplify the process.
2. Install Prerequisites: Ensure that you have the necessary tools installed, such as kubectl (the command-line tool for interacting with the cluster) and a container runtime.
3. Set Up the Control Plane: This can be done using tools like kubeadm, which helps bootstrap the cluster by initializing the Control Plane components.

4. Join Worker Nodes: Once the Control Plane is set up, you can join worker nodes to the cluster using the token generated during the initialization.
5. Deploy Applications: After the cluster is up and running, you can deploy your applications using YAML configuration files that define the desired state of your Pods and Services..

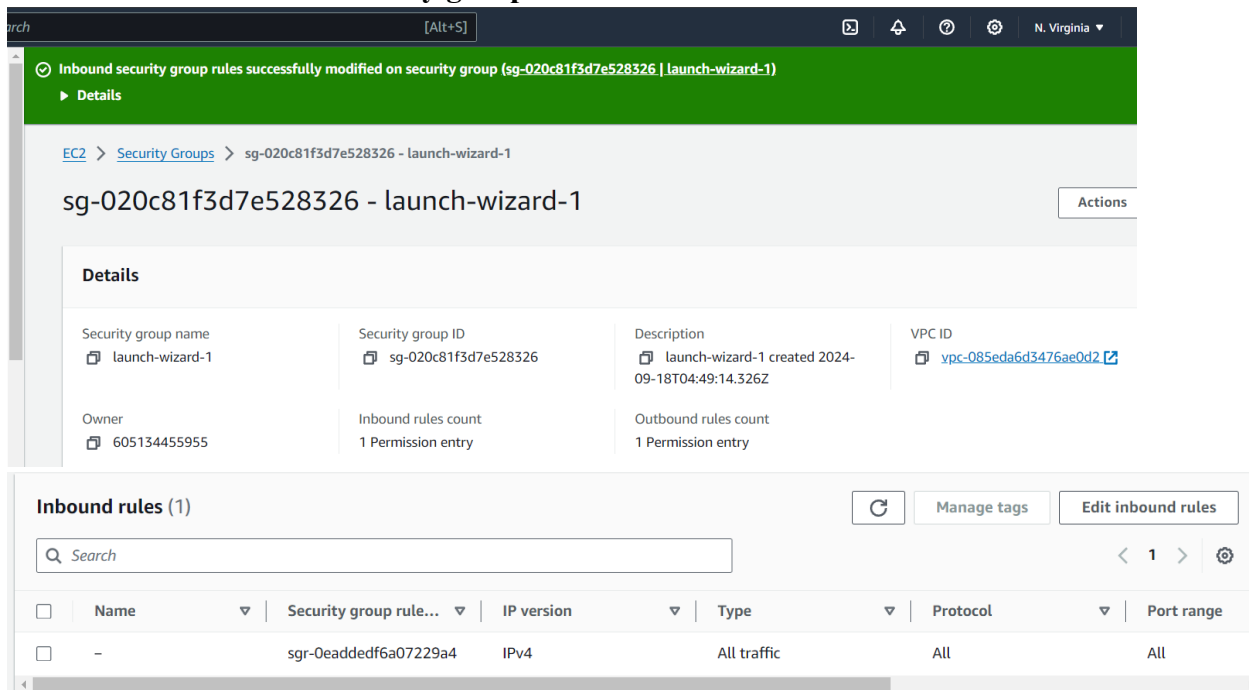
## Steps:

### 1. Creation of 2 EC2 Ubuntu Instances on AWS.



Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4
worker-node1	i-0526f6be95551ba74	Running	t2.micro	2/2 checks passed	View alarms	us-east-1a	ec2-3-81-7
master1	i-0aef82b0ccd222219	Running	t2.micro	2/2 checks passed	View alarms	us-east-1a	ec2-34-207

### 2. Edit inbound rules of security group 'launch-wizard-1' and set 'All Traffic'




sg-020c81f3d7e528326 - launch-wizard-1

Details			
Security group name launch-wizard-1	Security group ID sg-020c81f3d7e528326	Description launch-wizard-1 created 2024-09-18T04:49:14.326Z	VPC ID vpc-085eda6d3476ae0d2
Owner 605134455955	Inbound rules count 1 Permission entry	Outbound rules count 1 Permission entry	

Inbound rules (1)

Name	Security group rule...	IP version	Type	Protocol	Port range
-	sgr-0eaddedf6a07229a4	IPv4	All traffic	All	All

### 3. Set master and worker as hostname on respective servers



## 4.Installation of docker

### 4.1 - sudo apt-get update

```
aws | Services | Search | Docker | Notifications | Help | Settings | N. Virginia | ShravaniAnilPatil
ubuntu@master-node:~$ sudo apt-get update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease [126 kB]
Get:4 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
Get:5 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe amd64 Packages [15.0 MB]
Get:6 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/universe Translation
```

### 4.2 - sudo apt-get install docker.io

```
ubuntu@master-node:~$ sudo apt-get install docker.io
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  bridge-utils containerd dns-root-data dnsmasq-base pigz runc ubuntu-fan
Suggested packages:
  ifupdown aufs-tools cgroupfs-mount | cgroup-lite debootstrap
  docker-buildx docker-compose-v2 docker-doc rinse zfs-fuse | zfsutils
The following NEW packages will be installed:
  bridge-utils containerd dns-root-data dnsmasq-base docker.io pigz runc
  ubuntu-fan
0 upgraded, 8 newly installed, 0 to remove and 133 not upgraded.
Need to get 76.8 MB of archives.
```

### 4.3 – sudo systemctl enable docker

### 4.4 – sudo systemctl status docker

```
aws | Services | Search | Docker | Notifications | Help | Settings | N. Virginia | ShravaniAnilPatil
ubuntu@master-node:~$ sudo systemctl enable docker
ubuntu@master-node:~$ sudo systemctl status docker
● docker.service - Docker Application Container Engine
   Loaded: loaded (/usr/lib/systemd/system/docker.service; enabled; preset
   Active: active (running) since Wed 2024-09-18 05:07:22 UTC; 2min 35s ago
   TriggeredBy: ● docker.socket
     Docs: https://docs.docker.com
    Main PID: 2659 (dockerd)
      Tasks: 8
     Memory: 31.3M (peak: 33.2M)
        CPU: 293ms
     CGroup: /system.slice/docker.service
             └─2659 /usr/bin/dockerd -H fd:// --containerd=/run/containerd/c
```

## 5- Installation of Kubernetes-

### 5.1 sudo apt-get update

### 5.2 install ca certificate

```
aws | Services | Search | [Alt+S] | N. Virginia | ShravaniAnilPatil |
ubuntu@master-node:~$ sudo apt-get update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 http://security.ubuntu.com/ubuntu noble-security InRelease
Reading package lists... Done
ubuntu@master-node:~$ sudo apt-get install -y apt-transport-https ca-certificates curl
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
ca-certificates is already the newest version (20240203).
ca-certificates set to manually installed.
The following additional packages will be installed:
```

### 5.3 Download the public signing key for the Kubernetes package repositories.

```
aws | Services | Search | [Alt+S] | N. Virginia | ShravaniAnilPatil |
ubuntu@master-node:~$ sudo curl -fsSLo /usr/share/keyrings/kubernetes-archive-keyring.gpg https://packages.cloud.google.com/apt/doc/apt-key
ubuntu@master-node:~$
```

### 5.4 Add the appropriate Kubernetes apt repository

```
aws | Services | Search | [Alt+S] | N. Virginia | ShravaniAnilPatil |
ubuntu@worker1:~$ echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:stable/v1.31/deb/ /' | sudo tee /etc/apt/sources.l
ist.d/kubernetes.list
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:stable/v1.31/deb/ /
ubuntu@worker1:~$ sudo apt-get update
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Hit:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 http://security.ubuntu.com/ubuntu noble-security InRelease
Get:5 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:core:stable/v1.31/deb InRelease [1186 B]
Err:5 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:core:stable/v1.31/deb InRelease
The following signatures couldn't be verified because the public key is not available: NO_PUBKEY 234654DA9A296436
Reading package lists... Done
W: GPG error: https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:core:stable/v1.31/deb InRelease: The following signatures couldn't be verified bec
ause the public key is not available: NO_PUBKEY 234654DA9A296436
E: The repository 'https://pkgs.k8s.io/core:stable/v1.31/deb InRelease' is not signed.
N: Updating from such a repository can't be done securely, and is therefore disabled by default.
N: See apt-secure(8) manpage for repository creation and user configuration details.
```

### 5.5 sudo apt-get update

### 5.6 sudo apt-get install -y kubelet kubeadm kubectl

```
aws | Services | Search | [Alt+S] | N. Virginia | ShravaniAnilPatil
0% [Connecting to pkgs.k8s.io (34.107.204.206)] [Waiting for headers] [Connec
tp://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease
Hit:2 ht
0% [Connected to pkgs.k8s.io (34.107.204.206)] [Connecting to security.ubuntu
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:5 ht
tp://security.ubuntu.com/ubuntu noble-security InRelease
Get:4 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb InRelease [1186 B]
Get:6 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb Packages [4865 B]
Fetched 6091 B in 1s (7749 B/s)
Reading package lists... Done
ubuntu@master-node:~$ sudo apt-get install -y kubelet kubeadm kubectl
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following additional packages will be installed:
  contrack cri-tools kubernetec-cni
The following NEW packages will be installed:
  contrack cri-tools kubeadm kubectl kubelet kubernetec-cni
0 upgraded, 6 newly installed, 0 to remove and 130 not upgraded.
Need to get 87.4 MB of archives.
After this operation, 314 MB of additional disk space will be used.
Get:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 contrack amd64 1:1.4.8-1ubuntu1 [37.9 kB]
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble/main amd64 cri-tools amd64 1.31.1-1.1 [15.7 MB]
Get:3 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubeadm 1.31.1-1.1 [11.4 MB]
Get:4 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubectl 1.31.1-1.1 [11.2 MB]
Get:5 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubernetec-cni 1.5.1-1.1 [33.9 MB]
Get:6 https://prod-cdn.packages.k8s.io/repositories/iscv/kubernetes:/core:/stable:/v1.31/deb kubelet 1.31.1-1.1 [15.2 MB]
```

## 5.7 sudo apt-mark hold kubelet kubeadm kubectl

```
aws | Services | Search | N. Virginia | ShravaniAnilPatil
ubuntu@master-node:~$ sudo apt-mark hold kubelet kubeadm kubectl
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
ubuntu@master-node:~$
```

## 6. Kubernetes Deployment

### 6.1 sudo swapoff -a

### 6.2 Initialize Kubernetes on Master Node -

sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=all

```
aws | Services | Search | N. Virginia | ShravaniAnilPatil
table kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each
as root:

kubeadm join 172.31.33.243:6443 --token ltp4ao.gzdzdpkqly6bxvz \
--discovery-token-ca-cert-hash sha256:4423cf44f5102d477fa92160e76e03f
78f941d4e7a5836cd46bb14517dfab5ad
ubuntu@master-node:~$
```

i-Oaef82b0ccd222219 (master1)

PublicIPs: 34.207.105.187 PrivateIPs: 172.31.33.243

7.to create a directory for the cluster:

7.1mkdir -p \$HOME/.kube

7.2sudo cp -i /etc/kubernetes/admin.conf \$HOME/.kube/config

7.3sudo chown HOME/.kube/config

## 8. Deploy Pod Network to Cluster and Join Worker Node to Cluster

```
ubuntu@worker-node:~$ sudo kubeadm join 172.31.33.243:6443 --token ltp4ao.gzd2dpkqly6bxvzn --discovery-token-ca-cert-hash sha256:4423cf44f5102d477fa92160e76e03f78f941d4e7a5836cd46bb14517dfab5ad --ignore-preflight-errors=all
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 513.568999ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

ubuntu@worker-node:~$
```

Verify that everything is running and communicating:

8.1kubectl get pods --all-namespaces

8.2kubectl get nodes

```
ubuntu@master-node:~$ kubectl get pods --all-namespaces
NAME      NAMESPACE      NAME                                READY   STATUS    RESTARTS   AGE
kube-flannel  kube-flannel-ds-gx9xg  1/1      Running   0          14m
kube-flannel  kube-flannel-ds-tl79d  1/1      Running   0          6m2s
kube-system   coredns-7c65d6cfc9-nrns4  1/1      Running   0          23m
kube-system   coredns-7c65d6cfc9-pnh9p  1/1      Running   0          23m
kube-system   etcd-master-node      1/1      Running   0          23m
kube-system   kube-apiserver-master-node  1/1      Running   0          23m
kube-system   kube-controller-manager-master-node  1/1      Running   0          23m
kube-system   kube-proxy-8rz72      0/1      CrashLoopBackOff  7 (3m42s ago)  23m
kube-system   kube-proxy-w55hg      0/1      CrashLoopBackOff  4 (29s ago)    6m2s
kube-system   kube-scheduler-master-node  1/1      Running   0          23m

ubuntu@master-node:~$ kubectl get nodes
NAME             STATUS    ROLES    AGE   VERSION
master-node      Ready    control-plane  23m   v1.31.1
worker-node      Ready    <none>      6m22s v1.31.1

ubuntu@master-node:~$
```

